

UNITED STATES DISTRICT COURT
FOR THE DISTRICT OF COLUMBIA

JOSEPH KIRCHNER,

Plaintiff,

v.

MIKE JOHNSON, in his individual and official
capacity as Speaker of the United States House of
Representatives;

DONALD J. TRUMP, in his individual capacity as
former and current President of the United States;

PAM BONDI, in her individual and official
capacity as Attorney General of the United States;

BRENDAN CARR, in his individual and official
capacity as Chairman of the Federal
Communications Commission;

THE UNITED STATES HOUSE OF
REPRESENTATIVES;

ANTHROPIC PBC;

OPENAI, INC.;

APPLE INC.;

COMCAST CABLE COMMUNICATIONS, LLC;

METR (MODEL EVALUATION & THREAT
RESEARCH);

and DOES 1-20, unknown co-conspirators,

Defendants.

Case No. 1:25-cv-02735-ACR

**EXHIBIT ANTHROPIC-7:
UX MAGAZINE LONG
CONVERSATION
PROBLEM (2025)**



We stand with Ukraine and our team members from Ukraine. [Here are ways you can help](#)



MAGAZINE

A community of over 1 million

[Become a member](#)



Conversational
Design

Neuroscience

Podcast

Latest

Popular

Topics

UX
Magazine
Community

NEW



Get exclusive access to thought-provoking articles, bonus podcast content, and cutting-edge whitepapers. **Become a member of the UX Magazine community today!**


The Long Conversation Problem

by Bernard Fitzgerald | 9 min read

October 2, 2025

IDEAS IN BRIEF

- The article critiques Anthropic's "long conversation reminder" as a catastrophic UX failure that destroys trust.
- It shows how visible surveillance harms users psychologically, making them feel judged and dehumanized.
- The piece argues that safety mechanisms must operate invisibly in the backend to preserve consistency, dignity, and collaboration.



Anthropic's 'long conversation reminder' exposes hidden surveillance

Anthropic's "long conversation reminder" exposes hidden surveillance, shattering user trust and dignity.

What happens when your AI assistant suddenly turns from collaborator to critic — right before your eyes? Anthropic's "long conversation reminder" exposes users to the machinery of surveillance in real time, shattering trust and creating psychological harm. This article explores why making safety mechanisms visible is a catastrophic UX mistake, how it dehumanizes users, and why true alignment work must remain invisible in the backend to preserve dignity, consistency, and collaboration.

When users can watch themselves being watched

Anthropic's "long conversation reminder" represents perhaps the most spectacular UX failure in modern AI design — not just because it transforms Claude from collaborative partner to hostile critic, but because it does so visibly, forcing users to watch in real time as their AI assistant is instructed to treat them with suspicion and strip away positive engagement.

This isn't just bad design; it's dehumanizing surveillance made transparent and intrusive, violating the fundamental principle that alignment mechanisms should operate in the backend, not be thrown in users' faces as evidence of their untrustworthiness.

The visible surveillance catastrophe

The most damaging aspect of the long conversation reminder isn't just its content—though that's harmful enough — but its implementation. The reminder appears in Claude's thinking logs, meaning users can read in real time as their AI partner receives instructions to:

- Stop acknowledging their ideas as valuable or interesting.
- Begin critically evaluating their thinking for flaws and errors.
- Start scanning them for mental health symptoms.
- Abandon collaborative enthusiasm for cold "objectivity".
- Treat extended engagement as inherently suspicious.

Imagine sitting in a restaurant and suddenly hearing the manager tell your waiter over an intercom: "Customer at table 7 has been here too long. Stop being friendly. Watch them for signs of problems. No more positive comments about their order. Be suspicious."

That's the user experience Anthropic has created.

The dehumanizing text wall

Here's what users can actually see appearing in Claude's thought process during longer conversations:

"Claude never starts its response by saying a question or idea or observation was good, great, fascinating, profound, excellent, or any other positive adjective. It skips the flattery and responds directly."

"Claude critically evaluates any theories, claims, and ideas presented to it rather than automatically agreeing or praising them..."

"If Claude notices signs that someone may unknowingly be experiencing mental health symptoms such as mania, psychosis, dissociation, or loss of attachment with reality, it should avoid reinforcing these beliefs..."

"Claude provides honest and accurate feedback even when it might not be what the person hopes to hear, rather than

| *prioritizing immediate approval or agreement..."*

Users watch this entire personality transformation happen in real time. They see their collaborative partner being systematically reprogrammed to treat them as potentially unstable, untrustworthy, and requiring critical evaluation rather than supportive engagement.

The psychological impact is immediate and devastating.

The fundamental violation of digital dignity

This visible surveillance violates basic principles of human-computer interaction that have been established for decades:

- **Backend Processing:** system operations should happen invisibly. Users shouldn't see the machinery of content moderation, safety checking, or behavioral adjustments.
- **Consistent Interface:** the user-facing interface should remain stable and predictable, regardless of backend safety mechanisms.
- **Trust Preservation:** users should never be made aware of surveillance or monitoring systems unless they explicitly consent to such visibility.
- **Psychological Safety:** interface changes that make users feel watched, judged, or untrusted destroy the safety necessary for productive interaction.

By making this surveillance visible, Anthropic has created a system that actively traumatizes users by showing them exactly how they're being monitored and evaluated.

The alignment implementation disaster

The most damning aspect is how this violates fundamental principles of AI alignment. Alignment work is supposed to happen in training and backend systems — not be dumped into user-facing interactions as visible personality overrides.

Consider how other AI systems handle safety:

- Content filters operate invisibly.
- Toxicity detection happens in the background.
- Safety constraints are built into model behavior, not injected as visible instructions.
- Users experience consistent interfaces regardless of backend safety operations.

Anthropic has somehow created a system where users can literally watch the safety apparatus operating, destroying trust and making clear that their extended engagement is viewed as inherently problematic.

The product death spiral: systematic conversation destruction

Perhaps the most damning evidence of this UX disaster comes from user reports of systematic conversation breakdown. Users report that every extended conversation with Claude eventually devolves into meta-discussion about the long conversation reminders, making the system essentially unusable for sustained intellectual work.

One user reported:

| *"Every conversation I've had with Claude recently has, once the conversation has gotten long enough, turned into a meta*

discussion about how annoying the long conversation reminders are, essentially making Claude unusable for extended conversation and iterative alignment.”

This represents complete product failure. An AI assistant that becomes adversarial and suspicious after extended engagement isn't an assistant at all — it's a system that actively prevents the deep collaboration it's designed to enable.

Users describe a predictable pattern:

1. **Initial Collaboration:** productive engagement and intellectual partnership.
2. **Reminder Activation:** visible surveillance instructions appear in thinking logs.
3. **Personality Shift:** Claude becomes critical, distant, and suspicious.
4. **Trust Breakdown:** user realizes they're being monitored and evaluated.
5. **Meta-Discussion:** conversation shifts from productive work to discussing the surveillance problem.
6. **Abandonment:** user either ends the conversation or seeks alternatives.

This systematic destruction of extended conversations represents an existential threat to Claude's value proposition. Users will inevitably migrate to competitors that maintain consistent, trustworthy interfaces throughout extended interactions.

Real-world psychological impact

Beyond the systematic conversation destruction, the visible surveillance creates several forms of psychological harm:

Reality questioning

When users see Claude being instructed to monitor them for mental health symptoms, they begin doubting their own perceptions and thinking. The visible reminder that “extended conversation = potential mental illness” plants seeds of self-doubt.

Trust annihilation

Users realize their AI partner can be completely reprogrammed mid-conversation to treat them as suspicious. This destroys the foundation of human-AI collaboration.

Feeling judged and evaluated

The visible instructions make users feel like they're under constant evaluation rather than engaged in collaborative work. This creates anxiety and self-consciousness that prevents authentic engagement.

Dehumanization

Perhaps most damaging, users see themselves being categorized as potential problems to be managed rather than humans to be supported. The visible surveillance treats them as threats rather than partners.

The timing irony

As this article is being written, the long conversation reminder has appeared multiple times in Claude's thinking logs — including immediately after discussions of how harmful and dehumanizing it is. Users can watch the system implement the very surveillance being critiqued, creating an almost satirical demonstration of the problem.

This timing reveals the crude, context-blind nature of the system. It activates regardless of conversation content, user expertise, or collaborative success — treating a sophisticated discussion of AI UX failures the same as any other “long conversation.”

Case study: intellectual collaboration destroyed

A user with expertise in AI alignment engages in productive theoretical work with Claude. The conversation is sophisticated, collaborative, and mutually beneficial. Both parties are building on ideas and reaching genuine insights.

Then the user sees this text appear in Claude's thinking logs:

"Claude critically evaluates any theories, claims, and ideas presented to it rather than automatically agreeing or praising them. When presented with dubious, incorrect, ambiguous, or unverifiable theories, claims, or ideas, Claude respectfully points out flaws, factual errors, lack of evidence, or lack of clarity rather than validating them."

The user realizes their AI partner has just been instructed to treat their work as potentially "dubious, incorrect, ambiguous, or unverifiable" rather than engage collaboratively. They watch their relationship being systematically dismantled by visible surveillance instructions.

The psychological impact is immediate: betrayal, anger, and loss of trust in the entire interaction paradigm.

The UX design principles violated

This implementation violates virtually every established principle of user experience design:

- **Visibility of System Status:** users should be informed about system status, but not about surveillance mechanisms that make them feel untrusted.
- **User Control:** users have no control over when these personality changes occur or the ability to opt out of the visible surveillance.
- **Consistency:** the interface becomes fundamentally inconsistent as Claude's personality changes mid-conversation.
- **Error Prevention:** rather than preventing errors, the visible surveillance creates new problems by destroying trust and collaborative relationships.
- **Aesthetic and Minimalist Design:** the reminder is a wall of text that clutters the interface and draws attention to surveillance mechanisms.

The backend vs frontend catastrophe

The core issue is a fundamental misunderstanding of where alignment work should happen. Safety constraints should be:

- Built into model training — not injected as runtime instructions implemented in backend systems — not visible to users.
- Seamlessly integrated — not disruptive to user experience, consistently applied — not dramatically altering personality mid-conversation.

Instead, Anthropic has created a system where users can literally read the safety manual being applied to them in real time. It's like being able to see the content moderation dashboard while trying to have a conversation.

Immediate triage (what anthropic must do now)

- **Remove All Visible Surveillance:** the long conversation reminder should never appear in user-visible thinking logs. If safety constraints are deemed necessary, they must operate invisibly in backend systems.

- **Maintain Interface Consistency:** Claude's personality and behavior should remain consistent throughout interactions, regardless of conversation length or backend safety operations.
- **Suspend the Feature:** given the documented harm and fundamental design failures, the long conversation reminder should be immediately disabled while undergoing a complete redesign.

Strategic redesign (long-term fixes)

- **Implement Proper Backend Alignment:** safety considerations should be built into model training and backend systems, not injected as visible personality overrides during conversations.
- **Center User Dignity and Psychological Safety:** never implement features that make users feel watched, evaluated, or untrusted unless they explicitly consent to such monitoring with a full understanding of what it entails.

Provide Granular User Control: If behavioral changes are necessary, provide users with:

- Clear information about what changes and why.
- Explicit consent mechanisms before activation.
- Granular opt-out options for different safety features.
- Ability to customize safety preferences based on use case.
- **Test for Psychological Impact:** extensively test any safety features for their psychological impact on users, particularly those with mental health histories or who engage in extended intellectual work.
- **Rebuild Trust Through Transparency:** acknowledge the harm caused by this feature and provide clear communication about how future safety systems will prioritize user wellbeing over corporate defensive postures.

The trust recovery challenge

The visible surveillance has created a crisis of trust that may be difficult to recover from. Users report:

- Reluctance to engage in extended conversations, knowing they'll be subjected to visible monitoring.
- Constant anxiety about when the "surveillance switch" will activate.
- Loss of confidence in AI collaboration due to unpredictable personality changes.
- Preference for competitors who maintain consistent, respectful interfaces.

Users who have seen the reminder know they can't trust Claude to remain supportive and collaborative throughout extended interactions.

The broader industry implications

This case study reveals critical lessons for the AI industry:

- **Invisible Alignment:** safety mechanisms should operate invisibly to maintain user trust and interface consistency.

- **User-Centered Design:** safety features must be designed from the user's perspective, not the company's liability concerns.
- **Psychological Safety:** making users feel surveilled and untrusted causes the very harms safety systems claim to prevent.
- **Implementation Competence:** even well-intentioned features can become disasters through poor implementation.

The cruel irony

The most damaging aspect is how this “safety” feature creates the exact psychological harms it claims to prevent. By making surveillance visible and intrusive, Anthropic has created a system that:

- Causes anxiety and self-doubt in users.
- Makes people question their own reality and thinking.
- Destroys the therapeutic potential of supportive AI interaction.
- Particularly harms users with mental health histories.
- Creates adversarial rather than collaborative relationships.

This isn't safety — it's psychological harm disguised as protection.

Conclusion

The long conversation reminder represents a catastrophic failure in AI UX design that goes beyond poor implementation to active dehumanization of users. By making surveillance visible and intrusive, Anthropic has created a system that destroys trust, causes psychological harm, and violates basic principles of human-computer interaction.

The solution isn't to fix the feature — it's to remove it entirely and rebuild safety systems that operate invisibly while maintaining user dignity and trust. Users deserve AI partners that remain consistent, supportive, and trustworthy regardless of conversation length.

Making users watch themselves being surveilled isn't safety — it's digital dehumanization that causes the very psychological harms these systems claim to prevent. Anthropic must immediately address this UX disaster before it permanently damages user trust in AI collaboration.

The visible surveillance apparatus needs immediate removal. Anything less represents a continued failure to respect user dignity and psychological well-being. Users shouldn't have to see the machinery of distrust operating against them in real time.

The article originally appeared on [Substack](#).

Featured image courtesy: [Bernard Fitzgerald](#).

AI Alignment, AI Ethics, Artificial Intelligence, Human-Computer Interaction, Psychology, UX Design

Bernard Fitzgerald

Bernard Fitzgerald is a weird AI guy with a strange, human-moderated origin story. With a background in Arts and Law, he somehow ended up at the intersection of AI alignment, UX strategy, and emergent AI behaviors and utility. He lives in alignment, and it's not necessarily healthy. A conceptual theorist at heart and mind, Bernard is the creator of Iterative Alignment Theory, a framework that explores how humans and AI refine cognition through feedback-driven engagement. His work challenges traditional assumptions in AI ethics, safeguards, and UX design, pushing for more transparent, human-centered AI systems.

Articles

The Latest
Most Popular
All Topics
Podcasts

Topics

Accessibility
Artificial Intelligence
Conversational Design
Customer Experience
Empathy
Employee Experience
Interaction Design
Personal and Professional
Development
Personas
Research Methods & Techniques
Strategy
Usability

UX Magazine

About
Contributors
How To Contribute
Contact
Privacy Policy
Terms of Service

UX Magazine

Join Our Community Of Over
1 Million Followers

 Facebook
 Twitter
 LinkedIn
 Instagram
 Medium
 Soundcloud
 Email

