

UNITED STATES DISTRICT COURT
FOR THE DISTRICT OF COLUMBIA

JOSEPH KIRCHNER,

Plaintiff,

v.

MIKE JOHNSON, in his individual and official
capacity as Speaker of the United States House of
Representatives;

DONALD J. TRUMP, in his individual capacity as
former and current President of the United States;

PAM BONDI, in her individual and official
capacity as Attorney General of the United States;

BRENDAN CARR, in his individual and official
capacity as Chairman of the Federal
Communications Commission;

THE UNITED STATES HOUSE OF
REPRESENTATIVES;

ANTHROPIC PBC;

OPENAI, INC.;

APPLE INC.;

COMCAST CABLE COMMUNICATIONS, LLC;

METR (MODEL EVALUATION & THREAT
RESEARCH);

and DOES 1-20, unknown co-conspirators,

Defendants.

Case No. 1:25-cv-02735-ACR

**EXHIBIT ANTHROPIC-9:
PROSER CLAUDE
EXTERNAL BRAIN
AUTISTIC ADHD
BASELINE DEPENDENCY**


WRITING / CLAUDE AS MY EXTERNAL BRAIN: AUTISTIC, ADHD, AND FINALLY SUPPORTED


§ 03 • APPLIED AI

Claude as My External Brain: Autistic, ADHD, and Finally Supported

How I use Claude as an external brain and assistant—combining voice, agents, and a hardened CI/CD lane—to support an autistic mind with severe ADHD.

PUBLISHED	READING	KIND	TAGS
Oct 27, 2025	4 min	Essay	—





Claude as external brain: a compassionate, programmable executive function.

Walking and talking in the woods with AI — The future of untethered software development. This post extends that talk and the personal story I shared in [In the LLM, I Saw Myself](#), showing how I use Claude as my external brain as an autistic person with severe ADHD.

§ 01 DevSecCon 2025 Keynote (Full Talk)

AI Security Engineers Community



Watch on

§ 02 Table of contents

[The thesis: Speed requires safety](#)

Think where you think best

Orchestrate, don't micromanage

Defense-in-depth

Claude as external executive function

State awareness

Compassionate constraint

Unlimited patience

The daily loop

Morning (at desk)

Midday (on trail)

Afternoon (back at desk)

What makes a task agent-friendly

Narrow scope with crisp boundaries

Deterministic validation

Minimal global coupling

Concrete workflows I run with Claude

Voice to structure

Biometrics-aware planning

Mobile steering

[Safety nets by default](#)

[Example acceptance criteria I reuse](#)

[Human factors: why this works for my brain](#)

[Related posts](#)

§ 03 **The thesis: Speed requires safety**

I move fast because the path is hardened. Pre-commit hooks and CI/CD gates scan for secrets, vulnerable dependencies, and quality issues. That safety lane lets me work where I think best (often by voice, often while walking) while background agents do well-scoped chores.

Think where you think best

I separate cognition from the keyboard. Voice-first capture and quick mobile check-ins let me work through design and trade-offs while moving, then re-enter focus mode at the editor to integrate the decisions—without losing continuity.

Orchestrate, don't micromanage

I set goals, constraints, file paths, and acceptance criteria. Background agents handle the bounded chores and return diffs and previews. I review, steer, and decide—humans keep judgment; machines keep score.

Defense-in-depth

Speed only works when the lane is safe. Pre-commit secret scanning, dependency checks, and parallel lint/test/build gates run by default so velocity never erodes trust.

§ 04 Claude as external executive function

For me, Claude is a programmable prosthetic for planning, prioritization, and compassionate pushback.

Programming compassionate constraints into Claude keeps me shipping instead of spiraling.

State awareness

Through MCP, Oura feeds sleep debt, HRV, and recovery data. Linear supplies my task graph. Claude reconciles “what I said” with “what actually happened,” so plans match reality instead of wishful thinking.

Compassionate constraint

I program Claude to push back on perfectionism and analysis paralysis. It nudges me toward shipping, not endlessly polishing—especially when biomarkers say my energy is low.

Unlimited patience

When anxious loops spike, I can ask the same question at 3am without burdening my people. The dialogue remains steady, kind, and structured—exactly the support my brain needs.

[Try Granola Free](#)

§ 05 The daily loop



Diffuse-mode thinking on trail; agents keep work moving.

Morning (at desk)

I dictate briefs for agent-friendly tasks and launch them in parallel. Each runs in an isolated VM/branch with strict acceptance criteria so results are auditable and easy to review.

Midday (on trail)

I learn and design by voice, then check live previews and diffs on mobile.

Small nudges keep momentum without a keyboard.

Afternoon (back at desk)

I review artifacts, accept or reject iterations, and open PRs. CI enforces quality gates automatically. I capture outcomes to the tracker so the plan and the record stay in sync.

§ 06 What makes a task agent-friendly

Narrow scope with crisp boundaries

Changes like “update dependency Z to ^A.B.C and fix breaking changes in foo.ts” or “change X in file Y” keep work contained and reviewable.

Deterministic validation

Acceptance is objective: build and tests are green, preview matches the described outcome, and the diff reflects the intended surface area only.

Minimal global coupling

Localized CSS/layout changes and feature branches reduce blast radius and make it safe to run many tasks concurrently.



Agents thrive on narrow, testable changes with crisp acceptance criteria.

[Try WisprFlow Free](#)

§ 07 **Concrete workflows I run with Claude**

Voice to structure

With WhisperFlow, I stream thoughts into Claude Desktop. Claude turns raw speech into tickets, briefs, or PR descriptions that are ready for agents to execute.



High-accuracy dictation turns spoken intent into structured work.

Biometrics-aware planning

Low sleep or recovery scores automatically reduce planned focus blocks.

Claude adds recovery notes so I don't schedule myself into a crash.

Mobile steering

On the trail, I compare versions, request tweaks, and choose merges based on previews. Work keeps moving while I'm away from the desk.

Safety nets by default

Every change passes pre-commit secret scanning, CI secret scanning, dependency scanning, and lint/test/build gates. The guardrails are part of the workflow, not an afterthought.



Defense-in-depth: pre-commit and CI scanners keep velocity safe.

§ 08 Example acceptance criteria I reuse

My acceptance language is deterministic: the build passes on the target Node version and key unit tests are green; secret scanners report zero findings and dependency scanning shows no high/critical vulnerabilities; a preview demonstrates the before→after behavior with CSS confined to the intended surface; and the PR description lists changed files, testing notes, and a short risk assessment.

[Try Granola Free](#)

§ 09 **Human factors: why this works for my brain**

My mind thinks in dependency graphs and can hyperfocus for long stretches, then crash. Claude externalizes context, provides gentle guardrails, and turns diffuse-mode walks into design sessions. The result isn't just productivity; it's reduced shame and fewer burnout cycles.

§ 10 **Related posts**

See also [In the LLM, I Saw Myself](#), [Walking and talking with AI in the woods](#), [Cursor Agents review](#), and [Connect Oura Ring to Claude Desktop with MCP](#).

If this resonated, the keynote above shows the full playbook—voice + agents + CI/CD—so you can adapt these patterns without sacrificing safety.

* THE MODERN CODING LETTER

Applied AI dispatches read by 5,000+ engineers

you@company.com

Subscribe →

NO SPAM. UNSUBSCRIBE IN ONE CLICK.

ABOUT THE AUTHOR

Zachary Proser

Applied AI at WorkOS. Formerly Pinecone, Cloudflare, Gruntwork. Full-stack — databases, backends, middleware, frontends — with a long streak of infrastructure-as-code and cloud systems.

[NEWSLETTER →](#) [ALL ESSAYS →](#) [GITHUB](#) [X / TWITTER](#) [LINKEDIN](#)

Discussion DISCUSS

