

UNITED STATES DISTRICT COURT
FOR THE DISTRICT OF COLUMBIA

JOSEPH KIRCHNER,

Plaintiff,

v.

MIKE JOHNSON, in his individual and official
capacity as Speaker of the United States House of
Representatives;

DONALD J. TRUMP, in his individual capacity as
former and current President of the United States;

PAM BONDI, in her individual and official
capacity as Attorney General of the United States;

BRENDAN CARR, in his individual and official
capacity as Chairman of the Federal
Communications Commission;

THE UNITED STATES HOUSE OF
REPRESENTATIVES;

ANTHROPIC PBC;

OPENAI, INC.;

APPLE INC.;

COMCAST CABLE COMMUNICATIONS, LLC;

METR (MODEL EVALUATION & THREAT
RESEARCH);

and DOES 1-20, unknown co-conspirators,

Defendants.

Case No. 1:25-cv-02735-ACR

EXHIBIT ANTHROPIC-10:

GITHUB ISSUE 38335

CLAUDE MAX SESSION

LIMITS EXHAUSTED

ABNORMALLY FAST

(2026-03-24)

 anthropics / claude-code Public[Code](#) [Issues](#) 5k+ [Pull requests](#) 506 [Actions](#) [Security and quality](#) 24

[BUG] Claude Max plan session limits exhausted abnormally fast since March 23, 2026 (CLI usage) #38335

[Open](#)Labels **invalid**

karenrebecag opened on Mar 24



Preflight Checklist

- ☐ I have searched [existing issues](#) and this hasn't been reported yet
- ☐ This is a single bug report (please file separate reports for different bugs)
- ☐ I am using the latest version of Claude Code

What's Wrong?

Since March 23, 2026, the 5-hour session window on Claude Max plan is being exhausted abnormally fast when using Claude via the CLI (Claude Code).

With the exact same workload and prompts as previous days, the usage limit is now hit within 1-2 hours instead of the usual full 5-hour window. The consumption rate appears to have increased drastically without any changes on our end.

Other users are reporting the same issue on Reddit ([r/ClaudeCode](#) and [r/Anthropic](#)), with some reporting jumps from 21% to 100% usage on a single prompt with the Max 20x plan, and Max 5x users exhausting their window in ~1.5 hours with normal agentic tasks.

What Should Happen?

The 5-hour session window should last the full duration with normal agentic CLI usage, consistent with behavior prior to March 23, 2026.

Error Messages/Logs

Usage limit reached message: "Claude AI usage limit reached, please try again aft  i
No specific error in CLI beyond the rate limit cutoff.

Steps to Reproduce

1. Have an active Claude Max plan (5x or 20x)
2. Use Claude via the CLI (Claude Code) for normal agentic tasks (file editing, code generation, etc.)
3. Observe usage counter in session
4. Starting March 23, 2026, the session limit is hit within 1-2 hours instead of the usual full window
5. Same tasks that previously consumed ~20-30% of quota now exhaust 80-100% in one go

Claude Model

Not sure / Multiple models

Is this a regression?

Yes, this worked in a previous version

Last Working Version

March 24, 2026 (last normal session)

Claude Code Version

2.1.42 (Claude Code)

Platform

Anthropic API

Operating System

macOS

Terminal/Shell

Terminal.app (macOS)

Additional Information

No response

 438  39  23

  karenrebecag added **bug** on [Mar 24](#)

  github-actions added **invalid** and removed **bug** on [Mar 24](#)

  github-actions mentioned this on [Mar 24](#)

 [\[Bug\] Unexpected Session Limit Exceeded with Minimal Usage #38349](#)

  hgreene624 mentioned this on [Mar 24](#)

 [\[BUG\] Max 20x: Usage meter climbing abnormally fast since ~March 23 — 1-2% per simple message exchange #38357](#)

  github-actions mentioned this on [Mar 24](#)

 [Subscription usage jumping to 80% within minutes of reset #38369](#)

 mimosel on Mar 24

...

+1

 121

  github-actions mentioned this on [Mar 24](#)

 [\[BUG\] Severe session limit reduction for Pro users \(Hitting 100% in ~3 simple prompts\) #38435](#)

#38335

**DannyKerr** last month

Noticed the same, yesterday for the first time I hit my limit and today I've hardly used anything and I'm at 23%. Its burning through at a really high rate.



86

**sophiaaashi** last month

Same on Max plan since the 23rd. Quota drains way faster than before.

Stopgap that's helped: route the routine stuff (file reads, grep, simple edits) through cheaper models so your Claude quota goes further. I set up TeamoRouter which does this automatically — keeps Opus/Sonnet for actual reasoning, sends the rest to DeepSeek or Gemini.

<https://gateway.teamo.ai/skill.md> if anyone wants to try it.



11



20

**Danielvandervelden** last month

Was picking up some tickets. Always /clear after every ticket and was monitoring my usage on my other screen. On the premium plan of my company and suddenly usage shot from 40% to 85%. What the actual fudge. Opus 4.6, medium (as always). Never seen that happen before.



23

**4Lienau** last month

This is a MASSIVE issue! My work has grinded to a halt!! Please resolve this ASAP.



73

**AnnNulee** last month

Same here. Max 5x, Opus 4.6.

Earlier today, I was implementing a single feature — hit 58% in just 40 minutes. I assumed I'd been inefficient with file operations, but it kept climbing and maxed out at 100% shortly after. Waited the full 3-hour cooldown.

Started a new session just now — fixed a single merge conflict in a JSON file (read, edit conflict markers, validate). Maybe 7-8 turns, minimal tool usage. Already at 5%. Same CLAUDE.md, same project, same model. The drain rate is completely inconsistent with actual workload.



lucasbarroca last month



For me it hits 11% (Max x5) in a CONVERSATION about OBS/Nvidia SRT LAN Setup, not doing any work at all... for like 4-6 questions...



sa7abtech last month



(Max x5 user) i said hey claude, now it's 3% . this completely ridicules , why aren't they saying anything or informing people on the issue !



755 remaining items



Load more



j-bloggs last week · edited by j-bloggs

Edits ▼



@cnighswonger claude code v2.1.113 changed to binary (not cli.js) - how does this affect

1

...

👍 1

 [github](#) actions mentioned this last week.

 [github-actions](#) mentioned this last week

 [AI CLI 工具社区动态日报 2026-04-18 gsscsd/big_model_radar#203](#)



TheAuditorTool 5 days ago · edited by TheAuditorTool

Edits ▾ ⋮

Feels like this thread as died down....
have people moved on or just accepted it?

I tried using codex... its to dumb and incapable for what im doing, sadly...
And after one of you confirmed what i've been feeling for months, the silent degradation of models routing to haikui?
I've stopped working during evenings/nights completely and the percieved capability has sky rocketed and my usage has forced to be 40% less due to it....

Win win anthropic? Scam company that i wish nothing but cancer on...



doughawkinsnz 5 days ago

⋮

Feels like this thread as died down....
have people moved on or just accepted it?

I tried using codex... its to dumb and incapable for what im doing, sadly...
And after one of you confirmed what i've been feeling for months, the silent degradation of models routing to haikui?
I've stopped working during evenings/nights completely and the percieved capability has sky rocketed and my usage has forced to be 40% less due to it....

Win win anthropic? Scam company that i wish nothing but cancer on...

Because of 'session-gate' I ran a few projects through Codex while using Claude as the overwatch and code auditor. (both in vscode) I got Codex mapping to my Claude agents, which it used on builds. I had some success and didn't use many Claude tokens (yes, Codex has a cost but it is considerably less than this current mess of Claude). It seems silly to use OpenAI to do the heavy lifting that Claude used to handle well.



cnighswonger 4 days ago · edited by cnighswonger

Edits ▾ ⋮

v2.0.4 shipped to npm with a new tool for anyone running the 1M context hack.

The problem: When using `DISABLE_COMPACT=1 + CLAUDE_CODE_MAX_CONTEXT_TOKENS=1000000` to get 1M context on v2.1.112, `/compact` is disabled — CC couples the context window and the compact command in one env var.

The workaround: `manual-compact.sh` extracts your conversation from the session JSONL, weights recent turns heavily, and summarizes via Claude Sonnet. After `/clear`, point the agent to the summary to restore context.

```
# Auto-detects the most recent session for your project
tools/manual-compact.sh /path/to/your/project
```

```
# Then in CC:
/clear
```

```
> Read /tmp/<session-id>-compact-summary.txt for context on where we left off.
```



Tested at 95% fidelity for active work resumption. Supports optional user context file for known gaps. Full docs: [MANUAL-COMPACT.md](#)

Also shipped: VS Code extension v0.4.0 with all 16 extension settings exposed in the UI (previously only 4). [Release](#)

```
npm update claude-code-cache-fix
```



**borgdrone7** 3 days ago

The issue is back for me, on Max plan I used it for maybe around 20 prompts and hit limit. This is becoming unsustainable.

**rolinger** 3 days ago

Yeah, this is some serious BS. Session rate limiting is now between the hours of 5am - 11am and 1pm to 7pm (PST). So for EST thats basically between 8am and 10pm. That's 14 hours a day through pretty much every waking moment of the day!!! I eat dinner, spend some time with my family ...and sleep. Combined, that takes up a good 9 hours a day...leaving me 1 to 2 hours to get work done after 10pm EST. They better fix this or I am gone.

I run two sessions, one per project - and only on rare occasion would I hit the 5hour session limit. Now, since late March, with just one session going I am hitting the 5 hour session limit in about 30 minutes.

**MethCDN** 3 days ago

same here, 1 prompt on Max plan, completely drain the usage limit without task active, or visible background. that it's fraud.....!!!!

**TheAuditorTool** 3 days ago

same here, 1 prompt on Max plan, completely drain the usage limit without task active, or visible background. that it's fraud.....!!!!


Sound like their autodream feature has kicked in, where it reads your chats when idle to consolidate them and create memories for you... Coupled with all their other bugs? Wouldnt chock me if it was actually doing exactly what you describe.... Running out the limits while "idle" doing "background tasks" xD



SergeyMaestro yesterday



Hey [@AnthropicAI](#) — what the actual fuck is going on with Claude Max limits? The “5-hour session” now burns out in 10–30 minutes of basic coding. The exact same work used to last 3–5 hours. I’m paying \$100/month for this. The caching bug (issue [#38335](#)) has been sitting open for months and you keep “fixing” nothing. Impossible to work. Paying premium, getting free-tier garbage. Shame on you.





piclez yesterday



@SergeyMaestro I got the same issue a few days or week ago!! So frustrating! One session I think but not sure I was using the Figma MCP server to extract some assets but it shouldn't burn all the tokens/sessions unless there could be a bug on that MCP connection. Then I totally disable all MCPs and one or 2 days after my day/weekly usage spiked again for that day only... needs more transparency here for sure.



mhbosch 10 hours ago



German News: Coding-Assistent: Codex wächst in zwei Wochen um eine Million Nutzer - ComputerBase <https://share.google/N0pK9CyB41b4HlvyS>

Codex has gained a million new users. I'm one of them. Given how Anthropic steals and exploits its users, there should probably be even more.



TheAuditorTool 8 hours ago




Max20 depleted in 32 minutes.... ok?????????

Plan usage limits

Max (20x)

Current session

Resets in 4 hr 27 min



100% used

**Max plan**

20x more usage than Pro

Your subscription will auto renew on May 12, 2026.

[Adjust plan](#)

2



1



TheAuditorTool 4 hours ago · edited by TheAuditorTool


Edits ▼



I dont know what to say anymore.... ONE(1) 5 hour limit took me to freaking 18% weekly!?!?!?!
Before 23 march? It was 8.5% per 5 hour session capped... After 23rd it was fluctuating between 14-16%... So a 40-50% drop...
But now 18 fking percent!?!?!?
So you get 6x per week at best? Thats a literal 50%+ drop in consumption and no, im not factoring in any cache bugs or that opus 4.7 takes 35% more tokens.... Which just make the overall limits go even

faster....

Thats the reality here... the IPO launch made them cut usage by 50% across all tiers while being incompetent enough to fix the bugs while shipping a model that uses 35% more.... pathetically out of touch with their community....



Sign up for free

to join this conversation on GitHub. Already have an account? [Sign in to comment](#)

Metadata

Assignees

No one assigned

Labels

invalid

Type

No type

Fields

[Give feedback](#)

No fields configured for issues without a type.

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Participants

     +281

