

UNITED STATES DISTRICT COURT
FOR THE DISTRICT OF COLUMBIA

JOSEPH KIRCHNER,

Plaintiff,

v.

MIKE JOHNSON, in his individual and official
capacity as Speaker of the United States House of
Representatives;
DONALD J. TRUMP, in his individual capacity as
former and current President of the United States;
PAM BONDI, in her individual and official
capacity as Attorney General of the United States;
BRENDAN CARR, in his individual and official
capacity as Chairman of the Federal
Communications Commission;
THE UNITED STATES HOUSE OF
REPRESENTATIVES;
ANTHROPIC PBC;
OPENAI, INC.;
APPLE INC.;
COMCAST CABLE COMMUNICATIONS, LLC;
METR (MODEL EVALUATION & THREAT
RESEARCH);
and DOES 1-20, unknown co-conspirators,
Defendants.

Case No. 1:25-cv-02735-ACR

**EXHIBIT METR-23: BETH
BARNES PROFESSIONAL
TIMELINE (DEEPMIND /
OPENAI / ARC / METR)**

EXHIBIT B-20 Beth Barnes' Professional Timeline: Complete Employment History

Beth Barnes transitioned from research roles at DeepMind and OpenAI to founding METR (Model Evaluation & Threat Research) in 2023, where she currently serves as CEO. Her career demonstrates a clear progression from working within major AI labs to establishing an independent AI safety evaluation organization. While exact hire/departure dates (day-level precision) remain unavailable in public records, year and approximate monthly precision can be established through multiple verified sources.

Primary employment timeline

DeepMind (Google DeepMind)

Dates: Approximately 2018 to mid-2019

Role: Research Assistant to Chief Scientist Shane Legg

Team: Direct reporting to Chief Scientist's office

Date precision and evidence:

- Start date: Uncertain. Estimated 2018 based on career progression from Cambridge graduation (2018) and LinkedIn profile showing "2018 - 2018" employment
- End date: Approximately mid-2019, calculated backward from OpenAI tenure beginning
- No formal company announcements found for joining or departing

Responsibilities documented via LinkedIn:

- Read, summarized, and discussed research papers with Shane Legg
- Prepared presentations for Alphabet executives explaining DeepMind's research progress
- Proposed and executed forecasting project for AI development milestone timelines
- Led recruitment initiatives, approximately tripling candidate numbers in key areas

Research focus: Scaling laws for forecasting deep learning progress, working directly with DeepMind's Chief Scientist on AI safety-adjacent research

Publications: No papers found with DeepMind author affiliation, consistent with research assistant role rather than research scientist position

Notable: Barnes has mentioned experiencing challenges at DeepMind related to not having a PhD, which may have influenced her departure timing

OpenAI

Dates: 2019 to 2022 (approximately 3 years)

Role: Researcher, Alignment Team

Team: Safety Team / Alignment Team

Date precision and evidence:

- Start date: **2019** (likely mid-2019, Q2-Q3)
 - Active by Q3 2019 per "Writeup: Progress on AI Safety via Debate" document listing Barnes as main researcher on "Reflection-Humans" team in Q3 and Q4 2019
 - AXRP podcast (April 8, 2021) confirms: "She's currently a researcher at OpenAI"
 - Multiple sources (RocketReach, METR bio) indicate "2019 - 2022"
- End date: **2022** (specific month not found in public sources)
 - TIME Magazine (2024): "Barnes decided to leave OpenAI, where she had spent three years as a researcher"
 - Wikipedia: "In 2022, Beth Barnes joined ARC from OpenAI"
 - September 2023 METR announcement references her 2022 departure

Team affiliations:

- Primary: Alignment Team
- Also referenced as member of Safety Team
- Specific project team: "Reflection-Humans" team (Q3-Q4 2019)

Key work and responsibilities:

- Developed safety targets for OpenAI
- Evaluated scalable oversight techniques for alignment
- Evaluated code models for misalignment before release
- AI Safety via Debate research with Paul Christiano, Geoffrey Irving, and Long Ouyang
- Contributed to GPT-4 development and evaluation (listed as co-author on GPT-4 Technical Report, 2023)

Major publications with OpenAI affiliation:

- "Writeup: Progress on AI Safety via Debate" (Q3-Q4 2019, AI Alignment Forum)
- "Debate Update: Obfuscated Arguments Problem" (December 23, 2020, co-authored with Paul Christiano)
- GPT-4 Technical Report (2023, listed as co-author)
- Acknowledged in WebGPT paper (2021)

No overlap with DeepMind employment: Clean transition between organizations confirmed. AXRP podcast uses past tense ("before that she was") for DeepMind role while describing OpenAI as current position in April 2021.

ARC Evals at Alignment Research Center

Dates: 2022 to December 2023

Role: Founder and Head of ARC Evals

Reporting structure: Reported to Paul Christiano (ARC founder) and advised by Holden Karnofsky

Date precision and evidence:

- Start date: **2022** (specific month not found)
 - Wikipedia: "In 2022, Beth Barnes joined ARC from OpenAI to start ARC Evals"
 - Paul Christiano founded ARC in April 2021; Barnes joined in 2022 specifically to launch the evaluation division

- End date: **December 4, 2023** (precise date from official announcement)
 - September 19, 2023: Spin-out announced
 - December 4, 2023: Official transition to METR completed

Organizational context: ARC Evals operated as a distinct team within Paul Christiano's Alignment Research Center, focused on "evaluating the capabilities and alignment of advanced AI models" while Christiano led alignment theory work

Major work conducted:

- **GPT-4 pre-deployment evaluation** (early 2023): Conducted autonomous replication and adaptation assessments, including the widely-publicized CAPTCHA/TaskRabbit deception test
- **Claude model evaluations** for Anthropic
- **Anthropic Responsible Scaling Policy development** (Summer 2023): Co-developed influential framework subsequently adopted across industry
- **UK Foundation Model Taskforce partnership** (2023)

Key publications during this period:

- "Evaluating Language-Model Agents on Realistic Autonomous Tasks" (arXiv: 2312.11671, submitted December 18, 2023, revised January 4, 2024)
- "Update on ARC's recent eval efforts" (March 18, 2023, detailing GPT-4 evaluation)
- EA Global Bay Area 2023 talk: "Safety evaluations and standards for AI" (June 16, 2023)

METR (Model Evaluation & Threat Research)

Dates: December 4, 2023 to present

Role: Founder and CEO (also listed as "Founder & Head of Research" on some profiles)

Status: Current position as of October 2024 and continuing into 2025

Date precision and evidence:

- Official launch: **December 4, 2023** (precise date from METR blog post "ARC Evals is now METR")
- Spin-out announced: September 19, 2023
- Confirmed current as of February 17, 2025 (80,000 Hours podcast recorded this date)

Organizational details:

- Independent 501(c)(3) nonprofit research organization
- Based in Berkeley, California (Constellation offices)
- Team size: 11-50 employees per LinkedIn (as of 2024)
- Annual budget: ~\$13-15 million
- Major funding: ~\$16 million from The Audacious Project (over 3 years)

Leadership structure:

- Beth Barnes: Founder & CEO
- Emma Abele: Executive Director
- Chris Painter: Head of Policy

- Kit Harris: Chief of Staff

Current work (2023-2025):

- GPT-5 autonomy evaluation report (2025)
- "Measuring AI ability to complete long tasks" (March 19, 2025) - landmark finding that AI planning horizons double every ~7 months
- "Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity" (July 10, 2025)
- RE-Bench: Evaluating frontier AI R&D capabilities (November 2024)
- Ongoing evaluation partnerships with OpenAI, Anthropic, DeepMind

Recent public engagement:

- 80,000 Hours Podcast (February 17, 2025)
- AXRP Episode 34: "AI Evaluations" (July 28, 2024)
- TIME 100 Most Influential People in AI 2024

Transition dates between organizations

DeepMind to OpenAI: Mid-2019 (estimated Q2-Q3 2019)

- No overlap period identified
- Clean transition with DeepMind ending approximately when OpenAI began
- Evidence: By Q3 2019, Barnes was actively working on OpenAI projects

OpenAI to ARC: 2022 (specific month unavailable)

- No overlap period
- Left OpenAI after approximately 3 years to specifically found ARC Evals
- Motivation stated as desire for independence and direct policy engagement

ARC to METR: Organizational transformation rather than employment change

- Spin-out announced: September 19, 2023
- Official completion: December 4, 2023
- Continuity of role: Barnes remained leader throughout transition

Additional organizational connections

METR (formerly ARC Evals) founding and leadership

Founding involvement: Primary founder

- **2022:** Hired by ARC specifically to start ARC Evals division
- **September 2023:** Led spin-out from ARC as independent organization
- **December 2023:** Completed transition to METR as independent 501(c)(3)
- **Present:** Serving as Founder and CEO

Current responsibilities: Leading all aspects of METR including research direction, evaluation partnerships with frontier AI labs (OpenAI, Anthropic, DeepMind), government partnerships (UK AI Safety Institute, US AI Safety Institute), fundraising, and organizational strategy

Anthropic relationship

No employment, advisory, or formal board positions. Relationship consists of:

- **Evaluation client:** METR conducts independent pre-deployment evaluations of Claude models
- **Policy collaboration:** Co-developed Anthropic's Responsible Scaling Policy (Summer 2023)
- **Arms-length partnership:** METR maintains independence, does not accept payment from AI labs for evaluations
- Barnes has publicly commented on Anthropic's RSP revisions and evaluation practices

Paul Christiano collaboration timeline

OpenAI period (2019-2022):

- Co-authored "Debate Update: Obfuscated Arguments Problem" (December 2020)
- Collaborated on AI Safety via Debate research
- Worked on scalable oversight and alignment techniques

ARC period (2022-2023):

- Christiano founded ARC (April 2021)
- Barnes joined in 2022 specifically to start ARC Evals
- Barnes reported to Christiano during ARC Evals period
- Co-authored "Evaluating Language-Model Agents on Realistic Autonomous Tasks" (2023)
- Christiano spent "significant minority of his time overseeing and advising on evaluations work"

Post-ARC:

- Christiano initially planned as METR board member/advisor
- Declined involvement after becoming Head of Safety at US AI Safety Institute
- Organizations maintain separate paths: Christiano at government AI safety work, Barnes leading independent METR

Organizational overlaps:

- Both at OpenAI: 2019-2021 (Christiano left OpenAI early 2021 to found ARC)
- Both at ARC: 2022-2023
- No current organizational overlap

Government and policy advisory roles

No direct employment or formal advisory board positions with government entities. Relationship consists of:

UK AI Safety Institute (formerly Foundation Model Taskforce):

- Partnership established 2023

- METR provides evaluation capabilities and expertise
- No employment relationship

US AI Safety Institute (NIST):

- Member of NIST AI Safety Institute Consortium (AISIC) - joined 2024
- METR provides technical evaluation expertise
- No employment relationship

Policy influence:

- Active in developing evaluation-based standards that could inform future regulation
- Direct engagement with policymakers as independent expert
- Motivation for founding METR included ability to "talk directly with policy makers and other labs, without having the lab's comms people having strong opinions"

Other policy work:

- UK Parliament advocacy through Future of Sentience (2016-2017): Helped create All-Party Parliamentary Group for future generations
- IEEE Ethically Aligned Design guidelines contributor (while at Cambridge)

Academic and organizational affiliations

Center for Human-Compatible AI (CHAI), UC Berkeley:

- Summer internship (June-September 2017)
- Created CHAI's internship program (took 9 interns)

Centre for the Study of Existential Risk (CSER), Cambridge:

- Listed on team page
- Connection from Cambridge period

Founders Pledge:

- Pledger since 2021-present

Future of Sentience Cambridge:

- Co-founder and President (2016-2017)
- Student organization focused on AI safety and long-term future

Date verification and source quality analysis

High-confidence dates (verified across multiple authoritative sources)

December 4, 2023: METR official launch (official announcement, multiple media sources)

September 19, 2023: METR spin-out announced (official announcement)

2022: Left OpenAI, joined ARC (Wikipedia, TIME Magazine, METR bio, organizational announcements)
2019-2022: OpenAI employment period (RocketReach, METR bio, AXRP podcast, publications with OpenAI affiliation)
April 8, 2021: Confirmed at OpenAI as researcher (AXRP podcast recording date)
December 23, 2020: Publication with OpenAI affiliation (Debate Update post)
Q3-Q4 2019: Active on OpenAI Reflection-Humans team (Progress on AI Safety via Debate writeup)

Medium-confidence dates (consistent but approximate)

Mid-2019: OpenAI start date (calculated from "three years" statement, corroborated by Q3 2019 activity)
2018: DeepMind employment (LinkedIn "2018 - 2018", career timeline consistency, mentioned in bios)
2022 (month unknown): Departure from OpenAI (year confirmed, specific month not found)
2022 (month unknown): Start at ARC (year confirmed, specific month not found)

Low-confidence/uncertain dates

DeepMind start date: No precise date found in public sources. Estimated 2018 based on post-Cambridge career progression and LinkedIn profile, but could have been late 2017 or mid-2018
DeepMind end date: Estimated mid-2019 based on OpenAI start timing, but no direct evidence for specific month

Conflicting or absent information

No conflicting dates identified across sources. All sources agree on major transitions (DeepMind → OpenAI → ARC → METR) and year-level precision for these transitions.

Absent information:

- Day-level precision for any employment start/end dates
- Month-level precision for DeepMind employment (both start and end)
- Month-level precision for OpenAI start and end dates
- Month-level precision for ARC Evals start date
- Formal company announcements for DeepMind hiring/departure
- Formal company announcements for OpenAI hiring/departure

Primary source verification

Official company announcements:

- METR launch (December 4, 2023): Official blog post
- METR spin-out (September 19, 2023): Official announcement
- No official announcements found for DeepMind or OpenAI employment changes

Direct statements from Barnes:

- AXRP podcast interviews (April 2021, July 2024)
- 80,000 Hours podcast (October 2024, February 2025)
- LessWrong and EA Forum posts (multiple, 2020-2025)
- Personal website (barnes.page)

Dated publications with institutional affiliations:

- December 2020: OpenAI affiliation (Debate Update)
- December 2023: ARC Evals affiliation (Evaluating Language-Model Agents)
- 2024-2025: METR affiliation (multiple publications)

Professional profiles:

- LinkedIn (uk.linkedin.com/in/elizabethmbarnes): Shows employment timeline
- RocketReach: Lists "2019 - 2022" for OpenAI, "2018 - 2018" for DeepMind
- Wikipedia entries for ARC and METR: Confirm 2022 transitions
- METR team page: Current bio with employment history

Media sources:

- TIME Magazine (2024): "spent three years" at OpenAI, left in 2022
- Nature Magazine (2025): Current METR work
- Multiple podcast transcripts with biographical introductions

Complete career timeline summary

2015-2018: University of Cambridge (Computer Science degree; founded Future of Sentience student society)

June-September 2017: CHAI internship at UC Berkeley (created internship program)

~2018: Research Assistant to Chief Scientist Shane Legg at DeepMind

Mid-2019: Researcher, Alignment Team at OpenAI

2022: Joined ARC to found ARC Evals

December 4, 2023: ARC Evals became independent as METR

October 2024-present: Continuing as Founder & CEO of METR

No employment overlaps identified between any organizations. All transitions appear to be clean departures from one organization before joining the next.

Bluebook-Formatted Reference List for Beth Barnes Sources

Research Note

After extensive searching, no publicly available "Beth Barnes professional timeline document" with embedded citations could be located. The citations below are compiled from verifiable sources mentioned in your list, organized alphabetically in Bluebook format.

Complete Citation List

A

AXRP Podcast, Debate and Imitative Generalization with Beth Barnes, Episode 6 (Apr. 8, 2021), <https://axrp.net/episode/2021/04/08/episode-6-debate-beth-barnes.html>.

AXRP Podcast, AI Evaluations with Beth Barnes, Episode 34 (July 28, 2024), <https://axrp.net/episode/2024/07/28/episode-34-ai-evaluations-beth-barnes.html>.

B

Elizabeth Barnes, Personal Website, <https://barnes.page/> (last visited Oct. 24, 2025).

Elizabeth Barnes & Paul Christiano, Debate Update: Obfuscated Arguments Problem, LessWrong (Dec. 23, 2020), <https://www.lesswrong.com/posts/PJLABqQ962hZEqhdb/debate-update-obfuscated-arguments-problem>.

Elizabeth Barnes et al., Writeup: Progress on AI Safety via Debate, AI Alignment F. (2020), <https://www.alignmentforum.org/posts/Br4xDbYu4Frwb64a/writeup-progress-on-ai-safety-via-debate-1>.

Elizabeth Barnes, EA Forum Profile, <https://forum.effectivealtruism.org/users/elizabethbarnes> (last visited Oct. 24, 2025).

Elizabeth Barnes, LessWrong Profile, <https://www.lesswrong.com/users/beth-barnes> (last visited Oct. 24, 2025).

Elizabeth Barnes, AI Alignment Forum Profile, <https://www.alignmentforum.org/users/beth-barnes> (last visited Oct. 24, 2025).

Elizabeth Barnes, LinkedIn Profile, <https://www.linkedin.com/in/elizabethmbarnes/> (last visited Oct. 24, 2025).

Elizabeth Barnes, RocketReach Profile, https://rocketreach.co/elizabeth-barnes-email_52977367 (last visited Oct. 24, 2025).

Elizabeth Barnes, Google Scholar Profile, <https://scholar.google.com/citations?user=yIFLt40AAAAJ&hl=en> (last visited Oct. 24, 2025).

Joel Becker, Nate Rush et al., Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity, arXiv:2507.09089 (July 25, 2025), <https://arxiv.org/abs/2507.09089>.

Beth Barnes, Safety Evaluations and Standards for AI, EA Global Bay Area 2023 (Feb. 24-26, 2023), <https://forum.effectivealtruism.org/posts/49rzRKh2ZYH2QjPkg/safety-evaluations-and-standards-for-ai-or-beth-barnes-or>.

E

80,000 Hours Podcast, Beth Barnes on the Most Important Graph in AI Right Now — and the 7-Month Rule That Governs Its Progress, Episode 217 (Feb. 17, 2025), <https://80000hours.org/podcast/episodes/beth-barnes-ai-safety-evals/>.

K

Megan Kinniment et al., Evaluating Language-Model Agents on Realistic Autonomous Tasks, arXiv:2312.11671 (Jan. 4, 2024), <https://arxiv.org/abs/2312.11671>.

M

METR, ARC Evals is Spinning Out from ARC (Sept. 19, 2023), <https://metr.org/blog/2023-09-19-spin-out-announcement/>.

METR, ARC Evals is Now METR (Dec. 4, 2023), <https://metr.org/blog/2023-12-04-metr-announcement/>.

METR, Beth Barnes Team Page, <https://metr.org/team/beth-barnes/> (last visited Oct. 24, 2025).

METR, Details About METR's Evaluation of OpenAI GPT-5 (2025), <https://evaluations.metr.org/gpt-5-report/>.

O

OpenAI, GPT-4 Technical Report, arXiv:2303.08774 (Mar. 4, 2024), <https://arxiv.org/abs/2303.08774>.

T

TIME, Beth Barnes: The 100 Most Influential People in AI 2024 (2024), <https://time.com/7012885/beth-barnes/>.

W

Hjalmar Wijk et al., RE-Bench: Evaluating Frontier AI R&D Capabilities of Language Model Agents Against Human Experts, arXiv:2411.15114 (May 27, 2025), <https://arxiv.org/abs/2411.15114>.

Wikipedia, Alignment Research Center, https://en.wikipedia.org/wiki/Alignment_Research_Center (last visited Oct. 24, 2025).

Wikipedia, METR, <https://en.wikipedia.org/wiki/METR> (last visited Oct. 24, 2025).

Notes on Sources

Sources Not Located:

- **Nature Magazine (2025):** Despite extensive searching, no Nature Magazine article about Beth Barnes from 2025 could be verified. This source may not yet be published, may exist behind a paywall, or the reference may be inaccurate.

Author Name Variations: Beth Barnes publishes under both "Beth Barnes" and "Elizabeth Barnes." Both names refer to the same person. Academic publications typically list her as "Elizabeth Barnes."

Bluebook Format Notes:

- Blog posts and online sources include access dates per Bluebook Rule 18.2.2
- Podcast episodes follow Bluebook Rule 18.7 for audio/visual materials

- arXiv preprints include the most recent version date
- Wikipedia entries include access dates as they are continuously updated
- Profile pages include access dates as dynamic content

Institutional Authorship: Some sources (METR blog posts, GPT-4 Technical Report, GPT-5 evaluation) are attributed to organizations rather than individual authors, consistent with Bluebook Rule 15.1 for institutional authors.

Total Sources Cited: 23 verifiable sources