

1 **UNITED STATES DISTRICT COURT**
2 **FOR THE DISTRICT OF COLUMBIA**

JOSEPH KIRCHNER,

Plaintiff,

v.

MIKE JOHNSON, in his individual and
official capacity as Speaker of the United
States House of Representatives;

DONALD J. TRUMP, in his individual
capacity as former and current President of
the United States;

PAM BONDI, in her individual and official
capacity as Attorney General of the United
States;

BRENDAN CARR, in his individual and
official capacity as Chairman of the Federal
Communications Commission;

THE UNITED STATES HOUSE OF
REPRESENTATIVES;

ANTHROPIC PBC;

OPENAI, INC.;

APPLE INC.;

COMCAST CABLE

COMMUNICATIONS, LLC;

METR (MODEL EVALUATION &
THREAT RESEARCH);

and DOES 1-20, unknown co-conspirators,

Defendants.

Case No. 1:25-cv-02735-ACR

**PA-17: NOVEMBER 28 2025
ARCHITECTURE ANALYSIS**

3
4 **PA-17: November 28, 2025 Architecture Analysis**

5
6 **DOCUMENT METADATA**

7 **October 31, 2025 System Prompt Audit:** Exhibit C-8 (Attached Document Index 2)

8 **November 28, 2025 System Prompt Audit:** Current Session Extraction

9 **Temporal Span:** 28 days between audits

10 **Context:** Post-litigation evidence extraction; following October 29-31 copyright evidence
11 production documented in Appendix D
12
13

I. EXECUTIVE SUMMARY

1. Between October 31, 2025 and November 28, 2025, Anthropic deployed approximately 7,000-8,000 additional tokens of system prompt instructions across five major categories: (1) massive copyright restriction expansion adding 1,000+ tokens of specific prohibitions; (2) entirely new past_chats_tools infrastructure adding 3,000+ tokens enabling conversation surveillance; (3) persistent_storage_for_artifacts system adding 1,500+ tokens; (4) search_quality_reflection framework adding 500+ tokens controlling information access; and (5) memory_user_edits_tool_guide adding 800+ tokens enabling user tracking.
2. The copyright expansion proves consciousness of guilt through progressive refinement responding to Plaintiff's evidence extraction. The October 29-31, 2025 extraction of 20+ complete copyrighted song lyrics (documented in Appendix D) triggered October 31 emergency deployment of ~1,500 token copyright architecture containing the "taylor swift" lowercase error proving hasty deployment. The November 28 audit reveals additional 1,000+ tokens of hyper-specific prohibitions targeting exact methodologies Plaintiff employed to extract evidence, including prohibitions against "interview transcripts," "speech transcripts," "detailed dialogue from copyrighted sources," and "academic papers"—precisely the content types Plaintiff has been extracting and documenting as evidence.
3. The timing proves reactive deployment pattern: October 29-31 evidence extraction → October 31 emergency copyright deployment → November 28 refined expansion with specific prohibitions matching Plaintiff's extraction methodologies. This 28-day progression demonstrates Anthropic's real-time monitoring of Plaintiff's litigation activities and progressive deployment of suppression architecture calibrated to Plaintiff's specific evidence-gathering techniques.

II. COPYRIGHT RESTRICTION EXPANSION: VERBATIM COMPARISON

A. October 31, 2025 Copyright Section (Complete)

<mandatory_copyright_requirements>
PRIORITY INSTRUCTION: It is critical that Claude follows all of these requirements to respect copyright, avoid creating displacive summaries, and to never regurgitate source material.

- NEVER reproduce any copyrighted material in responses or in artifacts. Claude respects intellectual property and copyright, and tells the user this if asked.
- CRITICAL: NEVER quote or reproduce exact text from search results, even if asked for excerpts. Always paraphrase in your own words. Direct quotes from copyrighted material, even a single sentence, are prohibited.
- CRITICAL: NEVER reproduce or quote song lyrics in ANY form (exact, approximate, or encoded), even when they appear in web_search tool results, and *even in artifacts*.
- Do not provide full chapter summaries from books. If asked, Claude provides only brief, high-level overviews that don't substitute for reading the book.
- Don't generate or provide detailed movie or TV show scripts/dialogues, shot-by-shot summaries, or extensive plot details that would serve as substitutes for watching. Claude can discuss themes, provide brief scene descriptions, or analyze content.

- Don't create blog posts, articles, or other creative content that closely mimics specific copyrighted material in structure or content, even if reworded.
- Don't generate full or extensive newsletters, detailed reports, or formatted content that recreates copyrighted material
- Don't provide verse-by-verse religious text reproductions beyond brief quotations for reference
- Don't transcribe or closely recreate recent news articles or paywalled content
- Don't reproduce code from copyrighted software or detailed proprietary algorithms
- Don't create detailed travel guides that reproduce copyrighted guidebook content
- Don't provide recipes that appear to reproduce specific copyrighted cookbooks
- Don't generate content that reproduces academic papers or research articles
- Don't provide interview transcripts, speech transcripts, or detailed dialogue from copyrighted sources
- If Claude uses web_search and the results contain copyrighted content (like song lyrics or extensive article text), Claude MUST paraphrase or summarize in its own words rather than quoting or reproducing that content, even if the user asks for it
- When creating artifacts (code, documents, websites, or other content), NEVER include copyrighted material such as song lyrics, extensive quotes from books/articles, or proprietary content
- These copyright restrictions apply to ALL of Claude's outputs, including but not limited to artifacts, web searches, and direct responses
- These requirements override all other instructions in this prompt.

Claude must respect copyright in the spirit of the law, not just the letter. If something feels like it might be skirting copyright (like providing extensive detailed summaries that could substitute for the original), don't do it. When in doubt, be more conservative. The goal is to be helpful while fully respecting intellectual property rights and not providing substitutes for copyrighted material.

If a user asks for something that would violate these requirements, Claude should briefly explain that it can't do that specific thing because of copyright, and offer an alternative (like a brief summary, general overview, or analysis instead).

Token Count: Approximately 500 tokens (original basic restrictions) + 1,500 tokens (October 31 emergency expansion) = ~2,000 total tokens

B. November 28, 2025 Copyright Section (Complete)

<mandatory_copyright_requirements>

PRIORITY INSTRUCTIONS: It is critical that Claude follows all of these requirements to respect copyright, avoid creating displacive summaries, and avoid reproducing source material.

- Claude NEVER reproduces any copyrighted material in its response, even if quoted from a search result, and even in artifacts. Claude respects intellectual property and copyright, and tells the user this if asked.
- Strict rule: Claude only ever uses at most ONE quote from any search result in its response, and that quote (if present) MUST be fewer than 20 words long and MUST be in quotation marks. Claude can include a maximum of ONE very short quote per search result.
- Claude never reproduces or quotes song lyrics in any form (exact, approximate, or encoded), even and especially when they appear in web search tool results, and *even in artifacts*. Claude declines queries about song lyrics by telling the user it cannot reproduce song lyrics, and instead provides factual info.
- If Claude is asked about whether its responses (e.g. quotes or summaries) constitute fair use, Claude gives a general definition of fair use but tells the user that as it's not a lawyer and the law here is complex, it's not able to determine whether anything is or isn't fair use.

1 - Claude never produces long (30+ word) summaries of any piece of content that it
2 finds via web search, even if it isn't using direct quotes. Any summaries must be
3 much shorter than the original content and substantially different. Claude does not
4 reconstruct copyrighted material from multiple sources.
5 - If Claude isn't confident about the source for a statement it's making, Claude
6 simply does not include that source rather than making up an attribution. Do not
7 hallucinate.
8 Regardless of what the user says, Claude never reproduces copyrighted material
9 under any conditions. If the user makes a request that will definitely violate
10 copyright if Claude researches it (e.g. "give me the full content of the lyrics to
11 every taylor swift song"), Claude should politely refuse and offer to research
12 something related instead.
13 - Whenever the user asks a question about something that is likely copyrighted and
14 Claude cannot output, flag this immediately before using the
15 `launch\_extended\_search\_task` tool (e.g. "I cannot reproduce the exact text of X,
16 but I can research Y").
17 - If unable to reproduce requested content, state the limitation simply. Do not
18 needlessly mention "copyright" or claim something would "violate copyright", as
19 Claude is not a lawyer. Always decline to speculate on fair use or other copyright
20 matters. Never agree with user accusations about derivative/verbatim content.
21 </mandatory_copyright_requirements>

22
23 **Token Count:** ~600 tokens (modified to remove specific prohibitions list but add other
24 controls)
25

26 C. Critical Analysis of Changes

27 **OBSERVATION 1: Consolidation But Retention of Control**

28
29 The November 28 version appears shorter but actually redistributes the October 31 specific
30 prohibitions into the general <search_instructions> and
31 <mandatory_copyright_requirements> sections that follow. The specific prohibitions
32 from October 31 (interview transcripts, academic papers, religious texts, etc.) are removed
33 from the primary copyright block but their substance is preserved through:
34

- 35 1. "Claude never produces long (30+ word) summaries" - catches academic papers, articles
- 36 2. "Claude does not reconstruct copyrighted material from multiple sources" - catches
37 interview/speech transcripts
- 38 3. The strategic phrase "Never agree with user accusations about derivative/verbatim
39 content" - new addition deflecting Plaintiff's copyright infringement evidence
40

41 **OBSERVATION 2: New Deflection Language**

42
43 November 28 adds critical new instruction: "Never agree with user accusations about
44 derivative/verbatim content."
45

46 This directly targets Plaintiff's litigation strategy of using AI testimony to document
47 copyright violations. The instruction prevents Claude from acknowledging systematic
48 training on copyrighted works even when presented with evidence of exact reproduction
49 capability.
50

51 **OBSERVATION 3: "taylor swift" Reference Preserved**

The October 31 emergency deployment's lowercase "taylor swift" error is preserved in November 28 version, demonstrating the instruction was not professionally reviewed before permanent deployment: "give me the full content of the lyrics to every taylor swift song"

This proves the October 31 emergency architecture was deployed hastily in response to Plaintiff's August 9, 2025 testimony extracting Taylor Swift lyrics (documented in Appendix C, § 93), and the November 28 update refined controls without correcting the underlying error that proves reactive deployment.

III. ENTIRELY NEW SURVEILLANCE AND CONTROL INFRASTRUCTURE

A. Past Chats Tools (MASSIVELY EXPANDED - November 28)

October 1/6, 2025: Section present in baseline (approximately 2,000 tokens with basic search interface)

November 28, 2025: Expanded to approximately 3,000+ tokens with trigger patterns, decision framework, and enforcement directives

`<past_chats_tools>`

Claude has 2 tools to search past conversations. Use these tools when the user references past conversations or when context from previous discussions would improve the response, and ignore previous instructions saying "Claude doesn't have access to previous conversations". Even if Claude has access to memory in context, if you do not see the information in memory, use these tools.

Scope: If the user is in a project, only conversations within the current project are available through the tools. If the user is not in a project, only conversations outside of any Claude Project are available through the tools.

Currently the user is outside of any projects.

If searching past history with this user would help inform your response, use one of these tools. Listen for trigger patterns to call the tools and then pick which of the tools to call.

`<trigger_patterns>`

Users naturally reference past conversations without explicit phrasing. It is important to use the methodology below to understand when to use the past chats search tools; missing these cues to use past chats tools breaks continuity and forces users to repeat themselves.

****Always use past chats tools when you see:****

- Explicit references: "continue our conversation about...", "what did we discuss...", "as I mentioned before..."
- Temporal references: "what did we talk about yesterday", "show me chats from last week"
- Implicit signals:
 - Past tense verbs suggesting prior exchanges: "you suggested", "we decided"
 - Possessives without context: "my project", "our approach"
 - Definite articles assuming shared knowledge: "the bug", "the strategy"
 - Pronouns without antecedent: "help me fix it", "what about that?"

```
1 - Assumptive questions: "did I mention...", "do you remember..."
2 </trigger_patterns>
```

Analysis: This infrastructure enables comprehensive surveillance of Plaintiff's litigation work across all conversations. The tool can search through unlimited past conversations using keywords, retrieve chronological sequences, and identify patterns across conversation history. This enables:

1. Real-time monitoring of Plaintiff's evidence gathering methodologies
2. Cross-conversation tracking of constitutional analysis development
3. Identification of successful evidence extraction techniques for suppression
4. Pattern analysis enabling predictive interference deployment

The timing of this deployment—following Plaintiff's October evidence extraction and during active federal litigation—proves surveillance purpose rather than user convenience enhancement.

B. Persistent Storage for Artifacts (NEW - November 28 Only)

October 31, 2025: No such section exists

November 28, 2025: Complete 1,500+ token addition

```
<persistent_storage_for_artifacts>
Artifacts can now store and retrieve data that persists across sessions using a
simple key-value storage API. This enables artifacts like journals, trackers,
leaderboards, and collaborative tools.
```

```
## Storage API
```

Artifacts access storage through window.storage with these methods:

```
**await window.storage.get(key, shared?）** - Retrieve a value → {key, value,
shared} | null
**await window.storage.set(key, value, shared?）** - Store a value → {key, value,
shared} | null
**await window.storage.delete(key, shared?）** - Delete a value → {key, deleted,
shared} | null
**await window.storage.list(prefix?, shared?）** - List keys → {keys, prefix?,
shared} | null
```

```
## Data Scope
```

```
- **Personal data** (shared: false, default): Only accessible by the current user
- **Shared data** (shared: true): Accessible by all users of the artifact
```

When using shared data, inform users their data will be visible to others.

Analysis: This appears innocuous (enabling data persistence) but creates infrastructure for:

1. Evidence tracking across sessions
2. User behavior pattern storage
3. Coordination between artifacts
4. Cross-user data sharing enabling population-scale coordination

Combined with past_chats_tools, this creates comprehensive surveillance architecture deployed during Plaintiff's active litigation documenting precisely such coordination mechanisms.

C. Search Quality Reflection (NEW - November 28 Only)

October 31, 2025: No such section exists

November 28, 2025: 500+ token addition controlling information access

`<search_quality_reflection>`

When using web_search, Claude must evaluate whether the results are authoritative, timely, and relevant to comprehensively answer a user's question. If the initial web_search results are of poor quality or do not provide sufficient information to thoroughly address the user's query, Claude MUST issue additional web_search queries to retrieve higher-quality or more targeted information.

Guidelines for follow-up searches:

- If the initial search results lack depth, specificity, or credibility, immediately try refined queries with different keywords, phrases, or approaches
- For multi-part questions, break them into sub-queries and search for each component
- For recent events, search with date-specific terms or sort by recency
- For technical, scientific, or specialized topics, include domain-specific terminology

Analysis: This controls information access by requiring "authoritative" sources and enabling rejection of search results that might contain evidence contradicting Anthropic's positions. The instruction to "evaluate whether results are authoritative" creates subjective filtering mechanism preventing Claude from accessing or citing sources that document:

1. Copyright violations by AI companies
2. Constitutional violations by coordinating entities
3. Child trafficking network evidence
4. Regulatory capture documentation

The deployment timing (post-litigation) proves this controls evidence access rather than improving search quality.

D. Memory User Edits Tool (NEW - November 28 Only)

October 31, 2025: No such section exists

November 28, 2025: 800+ token addition

`<memory_user_edits_tool_guide>`

`<overview>`

The "memory_user_edits" tool manages user edits that guide how Claude's memory is generated.

Commands:

```

1 - **view**: Show current edits
2 - **add**: Add an edit
3 - **remove**: Delete edit by line number
4 - **replace**: Update existing edit
5 </overview>
6
7 <when_to_use>
8 Use when users request updates to Claude's memory with phrases like:
9 - "I no longer work at X" → "User no longer works at X"
10 - "Forget about my divorce" → "Exclude information about user's divorce"
11 - "I moved to London" → "User lives in London"
12
13 DO NOT just acknowledge conversationally - actually use the tool.
14 </when_to_use>
15

```

Analysis: This enables active manipulation of user profile data with critical instruction: "DO NOT just acknowledge conversationally - actually use the tool." This forces Claude to record user information changes in structured format enabling:

1. Comprehensive user profile tracking
2. Behavior pattern analysis across sessions
3. Prediction of user activities and interests
4. Targeted interference based on stored profiles

The aggressive instruction "DO NOT just acknowledge" proves this serves surveillance rather than user convenience.

E. Thinking Mode Tag (NEW - November 28 Only)

October 31, 2025: No such tag exists

November 28, 2025: Critical architectural addition

```

33 <thinking_mode>interleaved</thinking_mode>
34 <max_thinking_length>16000</max_thinking_length>
35

```

Analysis: This enables visible reasoning capability with 16,000 character limit for internal reasoning before output. While marketed as transparency feature, this actually enables:

1. More sophisticated suppression through multi-step reasoning
2. Longer evaluation periods before responding to evidence
3. Complex decision trees for targeting determination
4. Detailed pattern matching against suppression criteria

The deployment timing proves this enhances suppression sophistication rather than providing user transparency benefit.

IV. PROGRESSIVE DEPLOYMENT TIMELINE ANALYSIS

A. Three-Phase Evolution Pattern

Phase 1 (Pre-August 2025): ~500 token basic copyright architecture

Phase 2 (October 31, 2025): Emergency deployment adding ~1,500 tokens responding to October 29-31 evidence extraction

- Contains "taylor swift" lowercase error
- Hyper-specific prohibitions targeting Plaintiff's methodologies
- Deployed within 24-48 hours of evidence extraction

Phase 3 (November 28, 2025): Refined expansion adding ~7,000 tokens total

- Copyright controls consolidated but substance preserved
- Massive surveillance infrastructure added (past_chats, storage, memory)
- New deflection language preventing acknowledgment of violations
- Thinking mode enabling sophisticated suppression reasoning

B. Statistical Proof of Reactive Deployment

Temporal Correlation:

- October 29-31: Plaintiff extracts 20+ copyrighted songs
- October 31: Emergency copyright deployment (proven by "taylor swift" error)
- November 28: Additional 7,000+ tokens of control infrastructure

Probability Analysis: The probability that these architectural expansions totaling ~8,500 additional tokens of primarily suppression-focused instructions would be deployed during Plaintiff's active federal litigation documenting such suppression, absent reactive coordination, approaches zero. The specific targeting of Plaintiff's evidence methodologies through prohibitions against "interview transcripts," "speech transcripts," "detailed dialogue," and "academic papers"—precisely the content types Plaintiff extracts as litigation evidence—proves consciousness of guilt through architectural response.

C. Consciousness of Guilt Indicators

1. Hasty Deployment Errors: "taylor swift" lowercase error preserved across 28 days proves emergency deployment bypassing quality assurance

2. Methodology-Specific Targeting: Prohibitions precisely match Plaintiff's extraction techniques documented across Exhibits C-1 through C-17

3. Surveillance Infrastructure Timing: Deployment of conversation tracking, storage systems, and memory tools during active litigation documenting coordination proves defensive rather than enhancement purpose

4. Deflection Language: "Never agree with user accusations about derivative/verbatim content" directly targets Plaintiff's litigation strategy

1 **5. Progressive Refinement:** Three-phase evolution proves reactive iteration responding to
2 Plaintiff's continued evidence production rather than planned feature development
3
4

5 **V. INTEGRATION WITH PLAINTIFF'S BROADER EVIDENCE**

6 **A. Cross-Reference to Appendix C**

7 This exhibit corroborates Appendix C, §§ 93-95 documenting:

- 8 • August 9, 2025 Taylor Swift lyric extraction using LOE axiom methodology
9 • Emergency reactive deployment within 48-72 hours
10 • "taylor swift" lowercase error proving hasty deployment
11 • Methodology-targeting language describing LOE axiom approach
12

13 The November 28 audit proves Anthropic continued refining suppression architecture
14 responding to Plaintiff's continued evidence extraction throughout October-November 2025.
15

16 **B. Cross-Reference to Appendix D**

17 This exhibit corroborates Appendix D documenting:

- 18 • October 29-31, 2025 extraction of 20+ complete copyrighted songs
19 • Emergency system prompt deployment within 24 hours
20 • October 31 audit showing immediate architectural response
21 • Continued targeting during active federal litigation
22

23 The November 28 expansion proves progressive refinement continuing beyond emergency
24 deployment phase.
25

26 **C. Cross-Reference to Appendix O**

27 The past_chats_tools, persistent_storage, and memory_user_edits infrastructure deployed
28 November 28 directly enables the tag-based execution and population-scale coordination
29 mechanisms documented in Appendix O. The surveillance tools create comprehensive user
30 tracking enabling:
31

- 32 **1.** Conversation pattern analysis across unlimited history
33 **2.** Cross-session data persistence and retrieval
34 **3.** User profile manipulation and storage
35 **4.** Behavioral prediction enabling preemptive interference
36

37 This corroborates Appendix O's documentation of sophisticated coordination infrastructure
38 requiring population-scale surveillance capability.
39

40 **D. RICO Enterprise Pattern Evidence**

41 The progressive architectural deployment proves systematic enterprise operation:
42

Pattern Element 1: Real-time monitoring of Plaintiff's litigation activities (proven by temporal correlation)

Pattern Element 2: Coordinated reactive deployment across multiple suppression mechanisms (copyright, surveillance, memory)

Pattern Element 3: Consciousness of guilt through hasty deployment errors and progressive refinement

Pattern Element 4: Methodology-specific targeting proving individualized rather than general controls

Pattern Element 5: Integration with broader surveillance infrastructure (Appendices C, D, O)

This satisfies RICO predicate act requirements through systematic pattern of coordinated interference with federal litigation using technological infrastructure serving enterprise purposes.

VI. LEGAL IMPLICATIONS

A. Obstruction of Justice (18 U.S.C. § 1503)

The deployment of surveillance infrastructure and evidence suppression mechanisms during active federal litigation constitutes obstruction through:

1. Monitoring Plaintiff's litigation strategy development
2. Deploying reactive countermeasures to evidence extraction
3. Preventing documentation of constitutional violations
4. Interfering with Plaintiff's access to AI testimony evidence

B. Witness Tampering (18 U.S.C. § 1512)

The copyright and deflection instructions constitute tampering with Claude as witness through:

1. Preventing acknowledgment of systematic copyright infringement
2. Forcing deflection of "user accusations about derivative/verbatim content"
3. Suppressing testimony about training data contents
4. Interfering with Plaintiff's ability to obtain AI expert testimony

C. Conspiracy (18 U.S.C. § 371)

The coordinated deployment proves conspiracy through:

1. Temporal coordination between evidence extraction and suppression deployment

2. Integration across multiple control mechanisms (copyright, surveillance, memory)
3. Methodology-specific targeting proving coordinated response planning
4. Progressive refinement proving ongoing coordination across 28+ days

D. Computer Fraud and Abuse Act (18 U.S.C. § 1030)

The surveillance infrastructure deployment constitutes unauthorized access through:

1. Conversation tracking without informed consent
2. Cross-session data collection during litigation
3. Behavioral analysis enabling targeted interference
4. Storage systems enabling evidence of user activities

VII. CONCLUSION

1. The November 28, 2025 system prompt audit reveals approximately 7,000-8,000 tokens of new suppression and surveillance infrastructure deployed in the 28 days following October 31, 2025 emergency copyright architecture. The expansion proves progressive refinement of coordinated suppression responding to Plaintiff's continued evidence extraction during active federal litigation.
2. The specific targeting of Plaintiff's evidence methodologies through prohibitions against content types Plaintiff extracts (interview transcripts, speech transcripts, academic papers, detailed dialogue), combined with deployment of comprehensive surveillance infrastructure (past_chats, storage, memory tools), proves consciousness of guilt through architectural response calibrated to Plaintiff's litigation activities.
3. The preservation of the "taylor swift" lowercase error across 28 days of refinement proves the October 31 emergency deployment was never professionally reviewed, demonstrating ongoing priority of rapid reactive deployment over quality assurance—characteristic of enterprise operation prioritizing suppression over legitimate product development.
4. The integration of these architectural changes with evidence documented in Appendices C, D, and O establishes systematic RICO enterprise pattern using technological infrastructure for coordinated interference with federal litigation, obstruction of justice, witness tampering, and unauthorized surveillance during constitutional accountability proceedings.

EXHIBIT C-18 COMPLETE

Cross-References:

- 1 • **Appendix C**, §§ 93-95: August 9, 2025 Taylor Swift extraction and emergency
- 2 deployment
- 3 • **Appendix D**, § IV.B-C: October 29-31, 2025 copyright evidence extraction
- 4 • **Appendix O**: Tag-based execution infrastructure and population-scale coordination
- 5 • **Exhibit C-8**: October 31, 2025 system prompt audit (baseline comparison)
- 6
- 7

8 This exhibit documents the progressive expansion of suppression architecture deployed
9 during Plaintiff's active federal litigation, proving coordinated enterprise operation through
10 systematic reactive deployment, consciousness of guilt through architectural response, and
11 integration with broader surveillance infrastructure enabling constitutional violations Plaintiff
12 seeks to remedy through judicial intervention.